


```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.linear_model import LogisticRegression
from sklearn import metrics
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.metrics import precision_recall_curve
from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import (
    accuracy_score, confusion_matrix, classification_report,
    roc_auc_score, roc_curve, auc,
    ConfusionMatrixDisplay, RocCurveDisplay
)

from statsmodels.stats.outliers_influence import variance_inflation_factor
from imblearn.over_sampling import SMOTE
```


```
df=pd.read_csv('logistic_regression.csv')
```

```
df.head(5)
```



	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_owr
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	MOR
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	
3	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	6 years	
4	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	9 years	MOR
5 rows × 27 columns									

```
df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 235886 entries, 0 to 235885
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt              235886 non-null float64
1   term                  235886 non-null object
2   int_rate              235886 non-null float64
3   installment           235886 non-null float64
4   grade                 235886 non-null object
5   sub_grade             235886 non-null object
6   emp_title             222209 non-null object
7   emp_length            224948 non-null object
8   home_ownership        235886 non-null object
9   annual_inc            235886 non-null float64
10  verification_status    235886 non-null object
11  issue_d               235886 non-null object
12  loan_status           235886 non-null object
13  purpose               235886 non-null object
14  title                 234848 non-null object
15  dti                   235886 non-null float64
16  earliest_cr_line      235886 non-null object
17  open_acc              235886 non-null float64
18  pub_rec               235886 non-null float64
19  revol_bal             235886 non-null float64
20  revol_util            235715 non-null float64
21  total_acc             235886 non-null float64
22  initial_list_status    235886 non-null object
23  application_type       235886 non-null object
24  mort_acc              213302 non-null float64
25  pub_rec_bankruptcies   235557 non-null float64
26  address               235885 non-null object
dtypes: float64(12), object(15)
memory usage: 48.6+ MB
```

```
df.describe()
```

	loan_amnt	int_rate	installment	annual_inc	dti	open_acc	pub_rec	revol_bal	revol_util	total_i
count	235886.000000	235886.000000	235886.000000	2.358860e+05	235886.000000	235886.000000	235886.000000	2.358860e+05	235715.000000	235886.000000
mean	14104.732053	13.643355	431.524698	7.427901e+04	17.325831	11.306932	0.179167	1.581559e+04	53.777674	25.4241
std	8354.907949	4.467399	250.662467	6.006704e+04	8.130635	5.136844	0.544709	2.045974e+04	24.502087	11.8981
min	500.000000	5.320000	16.250000	2.500000e+03	0.000000	0.000000	0.000000	0.000000e+00	0.000000	2.0000
25%	8000.000000	10.490000	250.330000	4.500000e+04	11.260000	8.000000	0.000000	6.011000e+03	35.800000	17.0000
50%	12000.000000	13.330000	375.370000	6.400000e+04	16.880000	10.000000	0.000000	1.116300e+04	54.800000	24.0000
75%	20000.000000	16.490000	567.010000	9.000000e+04	22.960000	14.000000	0.000000	1.960375e+04	72.900000	32.0000
max	40000.000000	30.990000	1533.810000	7.446395e+06	189.900000	90.000000	86.000000	1.743266e+06	892.300000	151.0000

There is significant difference found in the mean and median of the following attributes

- loan_amnt
- terms
- installment
- revol_bal etc. These attributes might contain outliers

```
#checking for Non-Numeric Columns
cat_col=[col for col in df.columns if df[col].dtype=='O']
cat_col
```

```
['term',
 'grade',
 'sub_grade',
 'emp_title',
 'emp_length',
 'home_ownership',
 'verification_status',
 'issue_d',
 'loan_status',
 'purpose',
 'title',
 'earliest_cr_line',
 'initial_list_status',
 'application_type',
 'address']
```

```
#Number of Unique values from all non_numeric columns
for col in cat_col:

    print(f"No. of Unique values {col}: {df[col].nunique()}")
```

```
No. of Unique values term: 2
No. of Unique values grade: 7
No. of Unique values sub_grade: 35
No. of Unique values emp_title: 111427
No. of Unique values emp_length: 11
No. of Unique values home_ownership: 6
No. of Unique values verification_status: 3
No. of Unique values issue_d: 115
No. of Unique values loan_status: 2
No. of Unique values purpose: 14
No. of Unique values title: 31403
No. of Unique values earliest_cr_line: 664
No. of Unique values initial_list_status: 2
No. of Unique values application_type: 3
No. of Unique values address: 234999
```

```
#convert string data types into datetime format
df["earliest_cr_line"]=pd.to_datetime(df["earliest_cr_line"])
df['issue_d']=pd.to_datetime(df['issue_d'])
```

```
<ipython-input-8-a65e86cf7087>:2: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To
df["earliest_cr_line"]=pd.to_datetime(df["earliest_cr_line"])
<ipython-input-8-a65e86cf7087>:3: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To
df['issue_d']=pd.to_datetime(df['issue_d'])
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 235886 entries, 0 to 235885
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt             235886 non-null float64
1   term                  235886 non-null object
2   int_rate              235886 non-null float64
3   installment           235886 non-null float64
4   grade                 235886 non-null object
```

```

5   sub_grade      235886 non-null object
6   emp_title      222209 non-null object
7   emp_length     224948 non-null object
8   home_ownership 235886 non-null object
9   annual_inc     235886 non-null float64
10  verification_status 235886 non-null object
11  issue_d        235886 non-null datetime64[ns]
12  loan_status    235886 non-null object
13  purpose        235886 non-null object
14  title          234848 non-null object
15  dti            235886 non-null float64
16  earliest_cr_line 235886 non-null datetime64[ns]
17  open_acc       235886 non-null float64
18  pub_rec        235886 non-null float64
19  revol_bal      235886 non-null float64
20  revol_util     235715 non-null float64
21  total_acc      235886 non-null float64
22  initial_list_status 235886 non-null object
23  application_type 235886 non-null object
24  mort_acc       213302 non-null float64
25  pub_rec_bankruptcies 235557 non-null float64
26  address        235885 non-null object
dtypes: datetime64[ns](2), float64(12), object(13)
memory usage: 48.6+ MB

```

df.dtypes

```

↗ loan_amnt      float64
  term          object
  int_rate      float64
  installment    float64
  grade         object
  sub_grade     object
  emp_title     object
  emp_length    object
  home_ownership object
  annual_inc    float64
  verification_status object
  issue_d       datetime64[ns]
  loan_status   object
  purpose       object
  title         object
  dti           float64
  earliest_cr_line datetime64[ns]
  open_acc      float64
  pub_rec       float64
  revol_bal     float64
  revol_util    float64
  total_acc     float64
  initial_list_status object
  application_type object
  mort_acc      float64
  pub_rec_bankruptcies float64
  address       object
dtype: object

```

df.duplicated().sum()

```

↗ 0

```

df.isnull().sum()

```

↗ loan_amnt      0
  term          0
  int_rate      0
  installment    0
  grade         0
  sub_grade     0
  emp_title     13677
  emp_length    10938
  home_ownership 0
  annual_inc    0
  verification_status 0
  issue_d       0
  loan_status   0
  purpose       0
  title         1038
  dti           0
  earliest_cr_line 0
  open_acc      0
  pub_rec       0
  revol_bal     0
  revol_util    171
  total_acc     0
  initial_list_status 0
  application_type 0
  mort_acc      22584
  pub_rec_bankruptcies 329
  address       1
dtype: int64

```

We have bunch of missing value attributes.

df.describe(include = 'object')




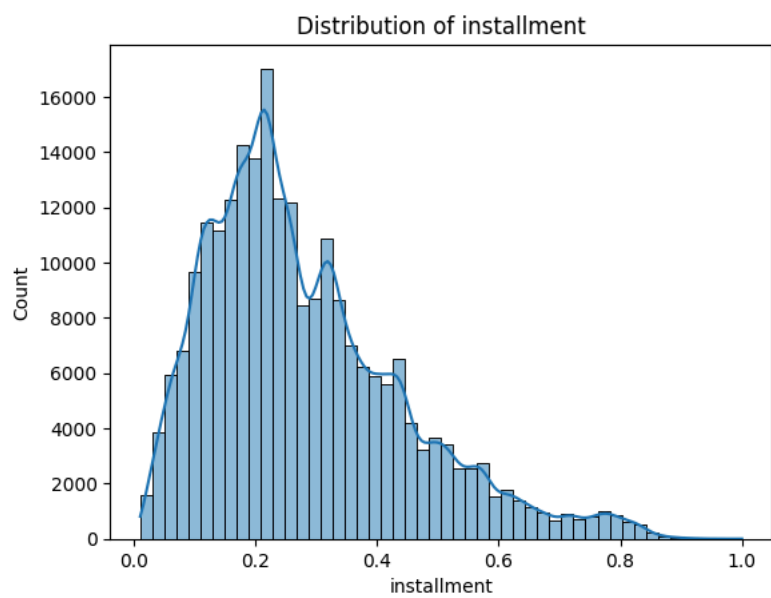

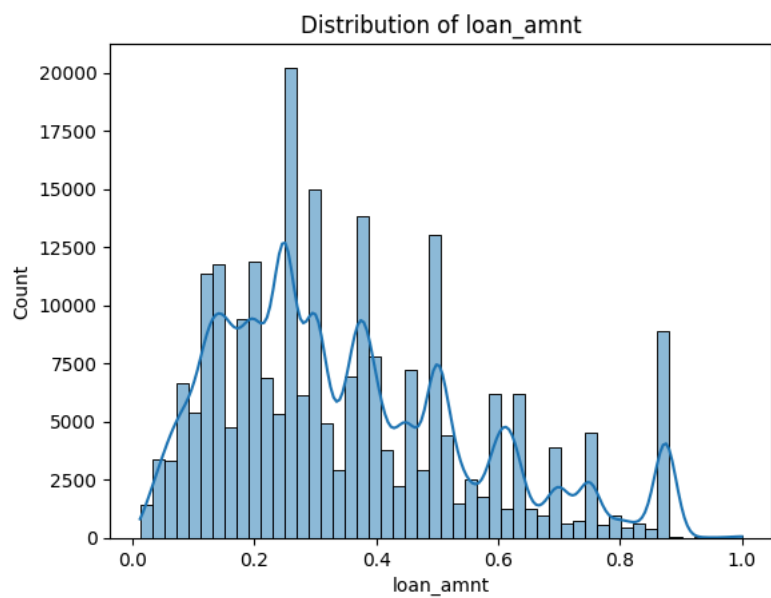
	term	grade	sub_grade	emp_title	emp_length	home_ownership	verification_status	loan_status		purpose	title	initial_1
count	235886	235886	235886	222209	224948	235886	235886	235886	235886	235886	234848	
unique	2	7	35	111427	11	6	3	2	14	31403		
top	36 months	B	B3	Teacher	10+ years	MORTGAGE	Verified	Fully Paid	debt_consolidation	Debt consolidation		
freq	179776	69059	15801	2607	75166	118154	82951	189716	139631	90798		


- Most of the loan disbursed for the 36 months period
- Most of the loan applicant have mortgage the home
- Majority of loans been fully paid off
- Majorily the loans been disbursed for the purpose of debt consolidation
- Most of the applicant is Individual

Univariate Analysis


num_vars = df.select_dtypes('float64').columns.tolist()

```
for i in num_vars:
    plt.title("Distribution of {}".format(i))
    sns.histplot(df[i]/df[i].max(), kde=True, bins=50)
    plt.show()
```







Distribution of dti




Distribution of open_acc

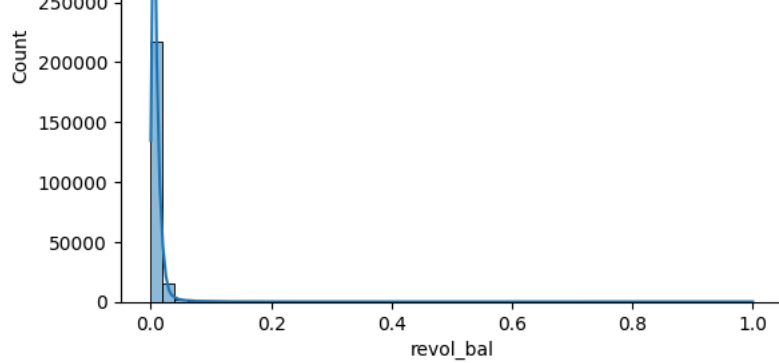


Distribution of pub_rec




Distribution of revol_bal







Distribution of revol_util




Distribution of total_acc



Distribution of mort_acc











Distribution of pub_rec_bankruptcies



- Most of the distribution is highly skewed towards the left side which tells us that they might contain outliers
- Almost all the continuous features have outliers present in the dataset.

```
cat_vars = ['home_ownership', 'verification_status', 'loan_status', 'application_type', 'grade', 'sub_grade', 'term']
for i in cat_vars:
    plt.figure(figsize=(8, 4))
    plt.title(f'Distribution of {i}')
    sns.countplot(data=df, x=i)
    plt.xticks(rotation = 45)
    plt.show()
```



- All the application type is Individual
- Most of the loan tenure is disbursed for 36 months
- The grade of majority of people those who have took the loan is 'B' and have subgrade 'B3'.
- So from that we can infer that people with grade 'B' and subgrade 'B3' are more likely to fully pay the loan.

Bivariate Analysis

```
plt.figure(figsize=(15,20))
```

```
plt.subplot(4,2,1)
sns.countplot(x='term',data=df,hue='loan_status')
```





```
plt.subplot(4,2,2)
sns.countplot(x='home_ownership',data=df,hue='loan_status')
```

```
plt.subplot(4,2,3)
sns.countplot(x='verification_status',data=df,hue='loan_status')
```

```
plt.subplot(4,2,4)
g=sns.countplot(x='purpose',data=df,hue='loan_status')
g.set_xticklabels(g.get_xticklabels(),rotation=90)
```

```
plt.show()
```

```
<ipython-input-17-cafc6a589ab6>:14: UserWarning: FixedFormatter should only be used together with FixedLocator
g.set_xticklabels(g.get_xticklabels(),rotation=90)
```




- Most of the people took loan term for 36 months and full paid on time
- Most of people have home ownership as mortgage and rent
- Most of the people took loan for debt consolidations



```
plt.figure(figsize=(15, 10))
```

```
plt.subplot(2, 2, 1)
grade = sorted(df.grade.unique().tolist())
sns.countplot(x='grade', data=df, hue='loan_status', order=grade)
```

```
plt.subplot(2, 2, 2)
sub_grade = sorted(df.sub_grade.unique().tolist())
g = sns.countplot(x='sub_grade', data=df, hue='loan_status', order=sub_grade)
g.set_xticklabels(g.get_xticklabels(), rotation=90)
```

```
plt.show()
```

 <ipython-input-18-4fc383fe14dc>:10: UserWarning: FixedFormatter should only be used together with FixedLocator
g.set_xticklabels(g.get_xticklabels(), rotation=90)




- The grade of majority of people those who have fully paid the loan is 'B' and have subgrade 'B3'.
- So from that we can infer that people with grade 'B' and subgrade 'B3' are more likely to fully pay the loan.



```
plt.figure(figsize=(15,12))

plt.subplot(2,2,1)
order = ['< 1 year', '1 year', '2 years', '3 years', '4 years', '5 years',
        '6 years', '7 years', '8 years', '9 years', '10+ years',]
g=sns.countplot(x='emp_length',data=df,hue='loan_status',order=order)
g.set_xticklabels(g.get_xticklabels(),rotation=90)
plt.title("Loan Years")

plt.subplot(2,2,2)
plt.barh(df.emp_title.value_counts()[:30].index,df.emp_title.value_counts()[:30])
plt.title("The most 30 jobs title afforded a laon")
plt.tight_layout()

plt.show()
```

 <ipython-input-19-d54df4ff7f6f>:7: UserWarning: FixedFormatter should only be used together with FixedLocator
g.set_xticklabels(g.get_xticklabels(),rotation=90)




- Person who employed for more than 10 years has successfully paid of the loan
- Manager and Teacher are the most afforded loan on titles

Correlatio Analysis

```
plt.figure(figsize=(10,8))
sns.heatmap(df.corr(numeric_only=True), cmap = 'crest', annot = True)
```

```
plt.show()
```



- We can see that almost perfect correlation between "loan_amnt" the "installment" feature.
- installment: The monthly payment owed by the borrower if the loan originates.
- loan_amnt: The listed amount of the loan applied for by the borrower. If at some point in time, the credit department reduces the loan amount, then it will be reflected in this value.

```
#drop column Installment
# df.drop(columns=['installment'],axis=1,inplace=True)
```

Feature Engineering

```
def pub_rec(number):
    if number == 0.0:
        return 0
    else:
        return 1

def mort_acc(number):
    if number == 0.0:
        return 0
    elif number >= 1.0:
        return 1
    else:
        return number

def pub_rec_bankruptcies(number):
    if number == 0.0:
        return 0
    elif number >= 1.0:
        return 1
    else:
        return number
```

```
df['pub_rec']=df.pub_rec.apply(pub_rec)
df['mort_acc']=df.mort_acc.apply(mort_acc)
```

```
df['mort_acc'] = df.mort_acc.apply(mort_acc)
df['pub_rec_bankruptcies'] = df.pub_rec_bankruptcies.apply(pub_rec_bankruptcies)
```

```
plt.figure(figsize=(12,20))

plt.subplot(6,2,1)
sns.countplot(x='pub_rec',data=df,hue='loan_status')

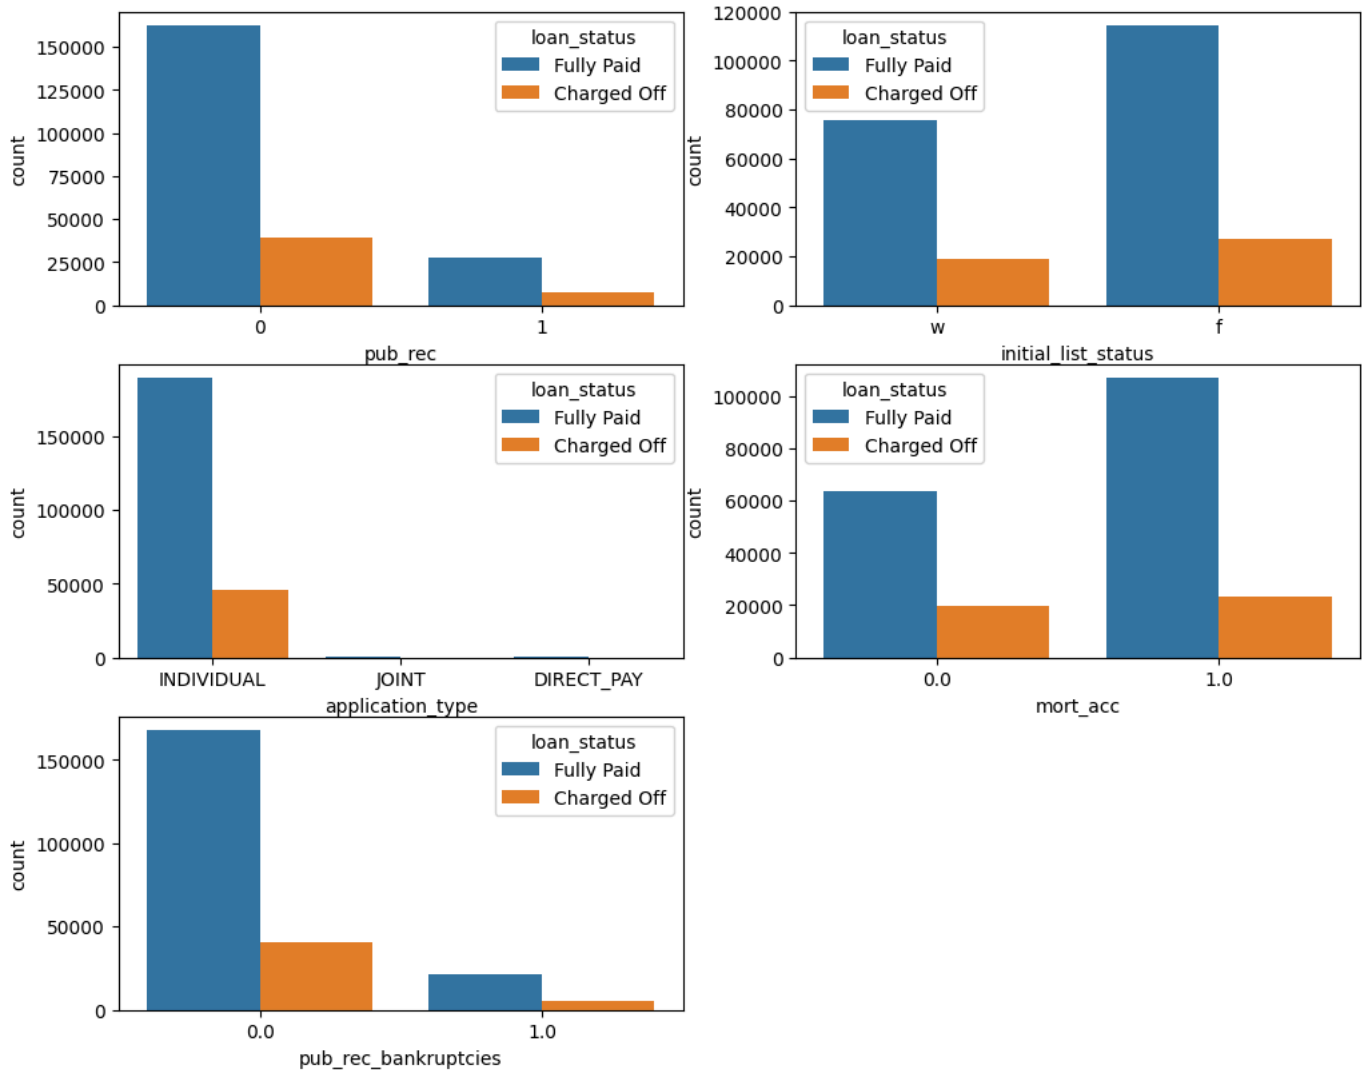
plt.subplot(6,2,2)
sns.countplot(x='initial_list_status',data=df,hue='loan_status')

plt.subplot(6,2,3)
sns.countplot(x='application_type',data=df,hue='loan_status')

plt.subplot(6,2,4)
sns.countplot(x='mort_acc',data=df,hue='loan_status')


plt.subplot(6,2,5)
sns.countplot(x='pub_rec_bankruptcies',data=df,hue='loan_status')

plt.show()
```



- As we can see that Most the loan disbursed to the people whose do not hold bankrupties record have successfully paid loan

```
df.describe()
```


	loan_amnt	int_rate	installment	annual_inc	issue_d	dti	earliest_cr_line	open_acc	pub_rec	
count	235886.000000	235886.000000	235886.000000	2.358860e+05	235886	235886.000000	235886	235886.000000	235886.000000	2
mean	14104.732053	13.643355	431.524698	7.427901e+04	2014-02-02 12:26:02.544618752	17.325831	1998-04-30 06:31:22.122720512	11.306932	0.146579	1
min	500.000000	5.320000	16.250000	2.500000e+03	2007-06-01 00:00:00	0.000000	1944-01-01 00:00:00	0.000000	0.000000	0
25%	8000.000000	10.490000	250.330000	4.500000e+04	2013-05-01 00:00:00	11.260000	1994-10-01 00:00:00	8.000000	0.000000	6
50%	12000.000000	13.330000	375.370000	6.400000e+04	2014-04-01 00:00:00	16.880000	1999-09-01 00:00:00	10.000000	0.000000	1
75%	20000.000000	16.490000	567.010000	9.000000e+04	2015-03-01 00:00:00	22.960000	2003-04-01 00:00:00	14.000000	0.000000	1
max	40000.000000	30.990000	1533.810000	7.446395e+06	2016-12-01 00:00:00	189.900000	2013-10-01 00:00:00	90.000000	1.000000	1
std	8354.907949	4.467399	250.662467	6.006704e+04	NaN	8.130635	NaN	5.136844	0.353687	2

▼ Default title text

```
# @title Default title text
```

```
df['total_acc'] = pd.to_numeric(df['total_acc'], errors='coerce')
```

```
df.head(5)
```

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ownership	annual_inc	...	open_acc	pub_rec	revol_bal
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	RENT	117000.0	...	16.0	0	36369.
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	MORTGAGE	65000.0	...	17.0	0	20131.
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	RENT	43057.0	...	13.0	0	11987.
3	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	6 years	RENT	54000.0	...	6.0	0	5472.
4	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	9 years	MORTGAGE	55000.0	...	13.0	0	24584.
5 rows × 27 columns														

```
df.isnull().sum()
```

	loan_amnt	0
	term	0
	int_rate	0
	installment	0
	grade	0
	sub_grade	0
	emp_title	13677
	emp_length	10938
	home_ownership	0
	annual_inc	0
	verification_status	0
	issue_d	0
	loan_status	0
	purpose	0
	title	1038
	dti	0
	earliest_cr_line	0
	open_acc	0
	pub_rec	0
	revol_bal	0
	revol_util	171
	total_acc	0
	initial_list_status	0
	application_type	0
	mort_acc	22584
	pub_rec_bankruptcies	329
	address	1
	dtype: int64	

```
# Dropping rows with null values
df.dropna(inplace=True)
```


```
# Remaining no. of rows
df.shape
```

```
↗ (199942, 27)
```


```
def box_plot(col):
    plt.figure(figsize=(8,5))
    sns.boxplot(x=df[col])
    plt.title('Boxplot for {}'.format(col))
    plt.show()

for col in num_vars:
    box_plot(col)
```



Boxplot for loan_amnt




Boxplot for int_rate




Boxplot for installment




Boxplot for annual_inc







Boxplot for dti



Boxplot for open_acc




Boxplot for pub_rec




Boxplot for revol_bal







Boxplot for revol_util




Boxplot for total_acc




Boxplot for mort_acc





Boxplot for pub_rec_bankruptcies




```
for col in num_vars:
    mean=df[col].mean()
    std=df[col].std()

    upper_limit=mean+3*std
    lower_limit=mean-3*std

    df=df[(df[col]<upper_limit) & (df[col]>lower_limit)]


df.shape
```

 (189438, 27)


```
def box_plot(col):
    plt.figure(figsize=(8,5))
    sns.boxplot(x=df[col])
    plt.title('Boxplot')
    plt.show()

for col in num_vars:
    box_plot(col)
```


Boxplot




Boxplot




Boxplot




Boxplot







Boxplot




Boxplot




Boxplot




Boxplot






Boxplot




Boxplot



Boxplot





```
# Converting term values to numerical val
term_values={' 36 months': 36, ' 60 months':60}
df['term'] = df.term.map(term_values)

# Mapping the target variable
df['loan_status']=df.loan_status.map({'Fully Paid':0, 'Charged Off':1})

# Initial List Status
df['initial_list_status'].unique()
np.array(['w', 'f'], dtype=object)
list_status = {'w': 0, 'f': 1}
df['initial_list_status'] = df.initial_list_status.map(list_status)

# Let's fetch ZIP from address and then drop the remaining details -
df['zip_code'] = df.address.apply(lambda x: x[-5:])
df['zip_code'].value_counts(normalize=True)*100
```

```
zip_code
70466    14.365650
30723    14.316557
22690    14.229986
48052    14.028864
29597    11.534117
05113    11.523559
00813    11.494526
11650     2.843147
93700     2.836812
86630     2.826782
Name: proportion, dtype: float64
```

```
# Dropping some variables which we can let go for now
df.drop(columns=['issue_d', 'emp_title', 'title', 'sub_grade',
                'address', 'earliest_cr_line', 'emp_length'],
        axis=1, inplace=True)
```


One hot encoding

```
dummies=['purpose', 'zip_code', 'grade', 'verification_status', 'application_type', 'home_ownership']
df=pd.get_dummies(df,columns=dummies,drop_first=True)
pd.set_option('display.max_columns',None)
pd.set_option('display.max_rows',None)
```

Data processing for modelling

```
from sklearn.model_selection import train_test_split
```


```
X=df.drop('loan_status',axis=1)
y=df['loan_status']
X_train, X_test, y_train, y_test =train_test_split(X,y,test_size=0.30,stratify=y,random_state=42)
print(X_train.shape)
print(X_test.shape)
```


 (132606, 50)
(56832, 50)


```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler()  
X_train = scaler.fit_transform(X_train)  
X_test = scaler.transform(X_test)
```

X_train

 array([[0.57104736, 0. , 0.27226581, ..., 0. , 0. ,
0.],
[0.66711141, 1. , 0.69173973, ..., 0. , 0. ,
1.],
[0.6490994 , 1. , 0.11859714, ..., 0. , 0. ,
0.],
...,
[0.18412275, 0. , 0.29303184, ..., 0. , 0. ,
1.],
[0.37358239, 0. , 0.13244116, ..., 0. , 1. ,
0.],
[0.37358239, 0. , 0.40470697, ..., 0. , 1. ,
0.]])



X_test

 array([[0.18679119, 0. , 0.28749423, ..., 0. , 0. ,
0.],
[0.4269513 , 0. , 0.26165205, ..., 0. , 1. ,
0.],
[0.37358239, 0. , 0.35394555, ..., 0. , 0. ,
1.],
...,
[0.72048032, 0. , 0.58467928, ..., 0. , 0. ,
0.],
[0.17344897, 0. , 0.26165205, ..., 0. , 0. ,
0.],
[0.40026684, 1. , 0.60775265, ..., 0. , 0. ,
0.]])


Model Building

```
from sklearn.linear_model import LogisticRegression
```

```
# Fit the LogisticRegression model on the preprocessed data  
logreg=LogisticRegression(max_iter=1000)  
logreg.fit(X_train,y_train)
```

  LogisticRegression
LogisticRegression(max_iter=1000)

```
pd.Series((zip(X.columns, logreg.coef_[0])))
```

 0 (loan_amnt, -0.04320443805806286)
1 (term, 0.5588719144348618)
2 (int_rate, -0.1909295408914164)
3 (installment, 0.5846791039181174)
4 (annual_inc, -1.0315790727485319)
5 (dti, 1.0103079618003479)
6 (open_acc, 0.7991834956794442)
7 (pub_rec, 0.13750498220070798)
8 (revol_bal, -0.39934958560523326)
9 (revol_util, 0.48983719714913054)
10 (total_acc, -0.6984359021795911)
11 (initial_list_status, 0.01798726724955526)
12 (mort_acc, -0.049546237270576714)
13 (pub_rec_bankruptcies, -0.16603105956947073)
14 (purpose_credit_card, 0.286135759732737)
15 (purpose_debt_consolidation, 0.3398576323440972)
16 (purpose_home_improvement, 0.3906316452021177)
17 (purpose_house, 0.24215443894893757)
18 (purpose_major_purchase, 0.3948450037229394)
19 (purpose_medical, 0.47399190105679645)
20 (purpose_moving, 0.3709326989026361)
21 (purpose_other, 0.3017251333476671)
22 (purpose_renewable_energy, 0.3401120689222868)
23 (purpose_small_business, 0.7671526300924224)
24 (purpose_vacation, 0.37073665091031516)
25 (purpose_wedding, -0.5012600259510492)
26 (zip_code_05113, -2.811547183676483)
27 (zip_code_11650, 11.563392239598716)
28 (zip_code_22690, 4.3333578375091015)

```
29         (zip_code_29597, -2.808033859868557)
30         (zip_code_30723, 4.374741129605297)
31         (zip_code_48052, 4.420405036331178)
32         (zip_code_70466, 4.3831363125349405)
33         (zip_code_86630, 11.547624020548648)
34         (zip_code_93700, 11.57884736333892)
35         (grade_B, 0.5495342892516659)
36         (grade_C, 1.0653674979435734)
37         (grade_D, 1.3950598653799136)
38         (grade_E, 1.6288226893735742)
39         (grade_F, 1.790221103052595)
40         (grade_G, 1.9088402431029736)
41     (verification_status_Source Verified, 0.190195...)
42     (verification_status_Verified, 0.0523628414920...)
43     (application_type_INDIVIDUAL, 0.7916536514798871)
44     (application_type_JOINT, -0.1452821665339001)
45     (home_ownership_MORTGAGE, -0.2585992735960012)
46     (home_ownership_NONE, -0.2344243559761461)
47     (home_ownership_OTHER, 0.48659722930795735)
48     (home_ownership_OWEN, -0.12966224859926567)
49     (home_ownership_RENT, 0.041645790966338914)
dtype: object
```

```
y_pred = logreg.predict(X_test)
print('Accuracy of Logistic Regression Classifier on test set: {:.3f}'.format(logreg.score(X_test, y_test)))
```

Accuracy of Logistic Regression Classifier on test set: 0.890

```
!pip install scikit-plot
!pip install scikitplot
!pip install scikitplot.metrics
```

Requirement already satisfied: scikit-plot in /usr/local/lib/python3.10/dist-packages (0.3.7)
Requirement already satisfied: matplotlib>=1.4.0 in /usr/local/lib/python3.10/dist-packages (from scikit-plot) (3.7.1)
Requirement already satisfied: scikit-learn>=0.18 in /usr/local/lib/python3.10/dist-packages (from scikit-plot) (1.2.2)
Requirement already satisfied: scipy>=0.9 in /usr/local/lib/python3.10/dist-packages (from scikit-plot) (1.11.4)
Requirement already satisfied: joblib>=0.10 in /usr/local/lib/python3.10/dist-packages (from scikit-plot) (1.4.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=1.4.0->scikit-plot) (1.2.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=1.4.0->scikit-plot) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=1.4.0->scikit-plot) (4.51.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=1.4.0->scikit-plot) (1.4.5)
Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=1.4.0->scikit-plot) (1.25.2)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=1.4.0->scikit-plot) (24.0)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=1.4.0->scikit-plot) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=1.4.0->scikit-plot) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=1.4.0->scikit-plot) (2.8.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.18->scikit-plot) (3.5.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib>=1.4.0->scikit-plot) (1.16.0)
ERROR: Could not find a version that satisfies the requirement scikitplot (from versions: none)
ERROR: No matching distribution found for scikitplot
ERROR: Could not find a version that satisfies the requirement scikitplot.metrics (from versions: none)
ERROR: No matching distribution found for scikitplot.metrics

Classification Report

```
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.99	0.94	45653
1	0.93	0.47	0.63	11179
accuracy			0.89	56832
macro avg	0.91	0.73	0.78	56832
weighted avg	0.89	0.89	0.87	56832

- Precision score and recall score for full paid status is almost same indicates that model is doing decent job which correctly classified the both of the scenarios
- Precision score for charged off status is more than recall score which is perfect

ROC / AUC

```
logit_roc_auc=roc_auc_score(y_test,logreg.predict(X_test))
fpr,tpr,thresholds=roc_curve(y_test,logreg.predict_proba(X_test)[:,-1])
plt.figure()
plt.plot(fpr,tpr,label='Logistic Regression (area = %0.2f)' % logit_roc_auc)
plt.plot([0,1],[0,1], 'r--')
```